

SAMEEKSHA KESHAV

sameeksha.keshav@gmail.com | +1 (984) 382-1285 | [linkedin.com/in/SameekshaKeshav](https://www.linkedin.com/in/SameekshaKeshav) | Raleigh, NC

SUMMARY

Master's student in Computer Science (Data Science track) at NC State University with 2+ years of experience spanning software engineering, data engineering, and applied AI/ML. Software Engineer at Microsoft's Azure DevDiv, building observability pipelines for Kubernetes-based distributed systems and large-scale ETL migrations. Currently building an agentic AI travel-planning chatbot at Tavant Technologies. Best Paper Award (IEEE CSITSS 2023). Strong in Python, C#, SQL, and Azure, with additional experience in NLP and LLM-based systems.

EDUCATION

North Carolina State University — M.S. Computer Science (Data Science Track) 2025 – Present

RV College of Engineering — B.Tech, Information Science (GPA: 9.70/10.0) 2020 – 2024

SKILLS

Languages: Python, C#, SQL, KQL, JavaScript, HTML/CSS

Data & AI/ML: ETL/ELT Pipelines, Data Modeling, TensorFlow, NLP, SBERT, Prompt Engineering, Agentic AI (Ollama), Model Evaluation

Cloud & Infra: Azure (Data Factory, Cosmos DB, Data Explorer, ML), AWS (EC2, S3, Lambda, Glue, Athena, EMR), Kubernetes, Git, CI/CD, .NET

Other: Pandas, NumPy, MySQL, Power BI, Excel, Jupyter Notebooks

WORK EXPERIENCE

Software Engineer, Microsoft India (R&D), Hyderabad Jul 2024 – Jul 2025

- Built and maintained logs and metrics pipelines for Kubernetes-based distributed systems on the Azure DevDiv observability team, supporting enterprise customers including Datadog, Dynatrace, New Relic, and Elastic.
- Engineered large-scale ETL migration pipelines (Cosmos MongoDB API → Cosmos DB Core API) using Azure Data Factory, KQL, C#, and .NET, resolving data discrepancies at enterprise scale.
- Automated manual workflows, improving distributed systems reliability and pipeline performance across the Azure DevDiv organization.

AIML Intern, NividApps Technology Solutions, Bangalore Feb 2024 – Jun 2024

- Developed a Vendor Recommendation System using AI/ML to streamline MSME procurement and improve vendor-selection accuracy within ERP workflows.

Software Engineer Intern, Microsoft India (R&D), Hyderabad May 2023 – Jun 2023

- Built an Azure Data Explorer dashboard analyzing service availability, reliability, and performance trends; deployed an Azure ML pipeline to predict delivery dates for internal orders.

Research Intern, Samsung Research Institute, Bangalore Jan 2023 – Aug 2023

- Performed big data analytics on game metadata using Python, t-SNE, and SBERT to power bundle recommendations, collaborating with a 3-person research team.

SELECTED PROJECTS

- [Inventory Management System](#):** Full relational database application (MySQL/Python) for a frozen-meals manufacturer, modeling multi-entity supply chain workflows on a normalized (BCNF/3NF) schema. Implemented stored procedures for transactional batch creation (FEFO lot selection, atomic inventory decrement, cost computation), triggers for lot-number generation and expired-lot prevention, compound ingredient cycle detection, and role-based access control for Manufacturer, Supplier, and Viewer roles.
- [Bat-Eye Bioacoustics Project \(NVIDIA-Sponsored\)](#):** Built TensorFlow deep learning pipelines for 9-class bat species classification from a 4GB ultrasonic recording dataset, with Bat Conservation India Trust and NVIDIA. Applied signal processing to extract features from raw ultrasonic audio; achieved 89.5% overall accuracy, with misclassification isolated to two acoustically similar species, showing strong generalization across the remaining 7.
- [State Employee Credit Card Analytics — Power BI \(PBIP\)](#):** Git-trackable Power BI project (PBIP/TMDL) analyzing 196,806 Delaware state P-Card transactions (~26MB). Built Power Query M transformations for data cleaning and feature engineering, statistical anomaly detection using per-category Z-scores (3 σ threshold) flagging 2,789 anomalous transactions (~1.4%), a star schema with a 291-row customer dimension, and 8 DAX measures across 2 report pages.
- [TimeTable Automation System](#):** Designed and deployed a production SQL-backed, Django-based timetable generation system in active use across 9 departments (~300–360 faculty) at RV College of Engineering. Automated constraint-solving for courses, faculty, and rooms with conflict-resolution and swap suggestions, reducing admin workload by 60% and cutting turnaround from 10 days to 4.
- [Student Performance Predictive Modeling](#):** Built a supervised ML system (Python/Django) on a 6,607×20 student performance dataset spanning academic, socio-economic, and behavioral factors. Led the regression component predicting continuous exam scores from study hours, attendance, sleep, and parental involvement, alongside classification tasks for extracurricular participation and score-improvement potential; deployed via a Django app with form-triggered inference.

PUBLICATIONS & HONORS

- Nayak, A. et al. (2024). "Popularity Estimation and New Bundle Generation using Content and Context based Embeddings." arXiv:2412.17310.
- N J, Prajwal et al. (2023). "Emotion Based Carnatic Music Recommendation System." IEEE CSITSS 2023 — Best Paper Award.
- Silver Medallist, B.Tech Computer Science (2nd Rank, Information Science, 2023); Microsoft Certified: Azure Data Fundamentals (DP-900) & Azure Fundamentals (AZ-900).